

Summary

My approach was to follow the recommended breakdown of the tasks as it seems to be a good way to break down the problem to intermediate tables that will help with the search. I only did the requirements, so no extra credit was approached.

The runtimes are as follows (includes printing to console, so I/O for that was added – see explanation in the Start Up Work section). However, it doesn't include the final saving of the file in the runtime calculation.

Time to perform Task 1: 0.988s

Time to perform Task 2: 105.214s

Time to perform Task 3: 102.798s

Time to perform Task 4: 104.624s

Total Runtime: 313.624s or 5.227 minutes

The program takes 3 inputs as required: <locations csv folder> <recordings folder> <output folder>. It assumes that the data files are the same as what we are given (2006.txt, 2007.txt, etc).

After Task 4, a .csv file is created with the results (though it matches what is in the report) into a folder called "result csv files" into the output folder (Spark seems to have issues where it will not allow me to use an existing folder as the direct directory). There will be a few files generated, but the one that is the resulting csv will look something like this: part-00000-04cf59f3-096a-4b92-a4e9-3a3eb490f6e9-c000.csv.

To run the program, simply run the .bat file with the correct parameters. Here is an example:

➤ `runproject.bat resources resources result`

Doing this will run the .jar file with the parameters. For this example, all my data files are in the resources folder and I want to output to a results folder. If the wrong folder is pointed, it will throw an error through Spark itself; if the wrong number of parameters are specified, there is a custom message outputted to the console.

Start Up Work

To start my project, I got Spark Core library version 3.0.0-preview2 (as any version earlier than 3 would not work with Windows due to the Hadoop library that Spark relies on had issues working with my Java version and other issues). In addition, I also downloaded Spark SQL of the same version to work with Spark as the tutorials from Apache Spark's main Github generally uses Spark SQL. To fix some other issues I had, I also needed to change the Jackson core and databind library to version 2.10.4. All this was done using Maven to update the dependencies.

Finally, I also added the environmental variable HADOOP_HOME pointing to winutils.exe for Hadoop. This is required for Windows computers to run the program. For example, mine points to: C:\hadoop where my winutils.exe is in a bin folder.

First, I needed to convert all the data to DataSet objects for manipulation; the WeatherStationLocations.csv was an easy convert as Spark allows for easy convert of a .csv. For the recordings text files, I needed to create JavaRDD objects and StructFields to represent the schema that I wanted for this data. Since there were a lot of fields we didn't need, I removed them early at this step. In addition, I added a calculated precipitation column for the different case letters there were and a month column where it gets the string of the month based on DATEMODA.

System.currentTimeMillis() was used to calculate Runtime. Because runtime is hard to calculate given the asynchronous methods for creating the Dataset, start time was placed prior to any variable creations for the task and end time was placed after printing the task. It is assumed that the program will not print rows for values that are not done, so this was the only way I could figure out how to ensure that the DataSet was fully created. However, there is a bit of overhead for printing, which I am unable to avoid due to this method of runtime calculation.

Task 1

To perform task 1, which is to get all the stations in the US and grouped them by state, the first thing I did was filter the results from the DataSet representing WeatherStationsLocations.csv by CTRY = "US" and removed all the results where the State filed is null (since I am not approaching the extra credit and averaging all entries missing the state wouldn't give you any meaningful data). I then grouped them by state and put all the USAF IDs into one column as an array.

The result was 53 rows for 53 "states" as the data seems to include VI (Virgin Islands), DC (District of Columbia), and PR (Puerto Rico). I did not remove these as I felt it was nice to have even if they were not real states.

Please see below for the results for this task:

```
+-----+-----+
|State|          USAFs|
+-----+-----+
| LA| [423630, 699560, ...|
| NM| [690014, 696334, ...|
| CA| [690020, 690070, ...|
| UT| [690044, 690064, ...|
| FL| [690090, 690524, ...|
| MI| [690110, 690260, ...|
| OK| [690130, 705022, ...|
| NC| [690138, 690200, ...|
| NV| [690170, 690174, ...|
| TX| [690190, 690190, ...|
| WA| [690230, 690240, ...|
| AZ| [690310, 696454, ...|
| OR| [690330, 720023, ...|
| AR| [690500, 720156, ...|
| AL| [691164, 720028, ...|
| MD| [691174, 720012, ...|
| AK| [691774, 692784, ...|
| IN| [692644, 720161, ...|
| GA| [692704, 720032, ...|
| MS| [695354, 699744, ...|
| VA| [695364, 699754, ...|
```

```

| CO| [698414, 720009, ...|
| IL| [699520, 720053, ...|
| HI| [699530, 720025, ...|
| SC| [699570, 720031, ...|
| SD| [711680, 720008, ...|
| MN| [720002, 720017, ...|
| MA| [720006, 720034, ...|
| OH| [720007, 720029, ...|
| IA| [720013, 720054, ...|
| NH| [720019, 720106, ...|
| RI| [720033, 722151, ...|
| ID| [720041, 720177, ...|
| ME| [720047, 720128, ...|
| KS| [720051, 720338, ...|
| MT| [720058, 720073, ...|
| PA| [720061, 720121, ...|
| NY| [720065, 720178, ...|
| WI| [720067, 720139, ...|
| MO| [720076, 720169, ...|
| CT| [720081, 720545, ...|
| ND| [720082, 720491, ...|
| NJ| [720116, 720179, ...|
| WV| [720160, 720160, ...|
| KY| [720163, 720353, ...|
| TN| [720168, 720171, ...|
| NE| [720308, 720308, ...|
| WY| [720341, 720345, ...|
| VT| [720492, 720493, ...|
| DE| [720495, 720898, ...|
| DC| [720618, 726402, ...|
| PR| [785140, 785145, ...|
| VI| [785430, 785510]|
+-----+

```

Time to perform Task 1: 0.988s

Task 2

To perform Task 2, I needed to manipulate the recordings Dataset. I did create a method to quickly create the Datasets separated for each recordings file, and eventually performed union on them to perform the task. Then, I needed to perform a join from Task 1 for the USAFs with the recordings' STN values to only have the recordings for the STN/USAFs that we care about. The union was performed with using Task 1 as the left DataSet since it is shorter than the recordings DataSet (Task 1 being only 53 states while the 2006 recordings text file alone has over 3 million lines). This would make it more efficient than if I were to do the join in the reverse direction.

After the join, I grouped each item by State and Month to perform average precipitation on the already calculated precipitation values (done at initialization of the DataSet), where the precipitation values were averaged when the state and month matched. Finally, I made a final DataSet where the month and average precipitation was aggregated into a single column and each state had a list of aggregated month/precipitation values (as an array). This was done using a special UDF (User Defined Function) in

Spark to combine the two column values together into a single value before I can use collect_list() to combine them into an array.

Please see below for the results for this task:

```
+-----+-----+
|State|      Precip List|
+-----+-----+
| AZ|[November 0.00893...|
| SC|[August 0.1457683...|
| LA|[November 0.08694...|
| MN|[November 0.01062...|
| NJ|[March 0.07900149...|
| DC|[January 0.0, Apr...|
| OR|[October 0.054720...|
| VA|[March 0.03249527...|
| RI|[August 0.0610893...|
| WY|[May 0.0636493374...|
| KY|[December 0.16628...|
| NH|[August 0.1361139...|
| MI|[April 0.04466275...|
| NV|[January 0.021177...|
| WI|[April 0.06165948...|
| ID|[April 0.04435056...|
| CA|[December 0.06309...|
| CT|[April 0.14679799...|
| NE|[February 0.01015...|
| MT|[May 0.0807792990...|
| NC|[August 0.0804419...|
| VT|[December 0.11962...|
| MD|[July 0.055024326...|
| DE|[October 0.103855...|
| MO|[March 0.15046365...|
| VI|[December 0.08630...|
| IL|[May 0.0871694619...|
| ME|[September 0.1328...|
| WA|[April 0.06064643...|
| ND|[January 0.014514...|
| MS|[November 0.10504...|
| AL|[March 0.12291313...|
| IN|[May 0.1100024021...|
| OH|[July 0.144883138...|
| TN|[April 0.16501446...|
| NM|[January 0.012061...|
| IA|[May 0.0583075817...|
| PA|[May 0.1073579339...|
| SD|[May 0.0886068256...|
| NY|[January 0.077787...|
| TX|[May 0.0775800039...|
| WV|[October 0.095235...|
| GA|[July 0.111734122...|
| MA|[August 0.1164332...|
| KS|[October 0.116547...|
| CO|[December 0.02631...|
```

```

| FL|[February 0.09714...|
| AK|[November 0.04172...|
| AR|[June 0.131780525...|
| OK|[December 0.03375...|
| PR|[June 0.035441176...|
| UT|[July 0.021499228...|
| HI|[April 0.04544023...|
+-----+

```

Time to perform Task 2: 105.214s

Task 3

To perform Task 3, I needed the DataSet I created for Task 2. The column with the month/average precipitation parsed using a special UDF where I looked at each month/precipitation value and found which was the lowest and highest month. Because I had to aggregate the month/precipitation into an array in Task 2, I had made them into Strings with a delimiter of a single space. Thus, I created an intermediate DataSet where I didn't parse the String and a DataSet using this DataSet object to split the month/precipitation column to corresponding columns for lowest month, lowest precipitation, highest month, and highest precipitation values.

Please see below for the results for this task:

```

+-----+-----+-----+-----+-----+
|State|Lowest Month|Lowest Precipitation|Highest Month|Highest Precipitation|
+-----+-----+-----+-----+-----+
| AZ|    November|0.008935691105371218|      July| 0.07065813761059134|
| SC|    November| 0.07339613086194678|    August| 0.14576830669591567|
| LA|    November| 0.0869480834546337|      July| 0.1697039196631033|
| MN|     January|0.003336326793773...|      June| 0.047394458223324575|
| NJ|    February| 0.07410393700787403|      July| 0.1785957240038873|
| DC|     January|                0.0|    January|                0.0|
| OR|       July| 0.00966687020829092|    November| 0.14075152653828096|
| VA|    February|0.024673968515683268|      June| 0.08022623892745996|
| RI|    February| 0.05318513451892387|    April| 0.08669625871908687|
| WY|    February|0.011197420634920623|      May| 0.06364933741080532|
| KY|    November| 0.09035292213048407|      July| 0.16976356251813182|
| NH|     March| 0.08385728710871762|      July| 0.23177604593929466|
| MI|    February| 0.03219802952210654|    September| 0.07008190283855044|
| NV|     August|0.009235637779941571|      July| 0.027827986142491327|
| WI|     January|0.017938713700338508|      June| 0.08139457959880501|
| ID|     August| 0.01603681242685391|    November| 0.06329878524462283|
| CA|       July|0.002754643909795044|    February| 0.08149602463781755|
| CT|     March| 0.07905631114292533|      June| 0.1607084785133567|
| NE|     January|0.007690058479532156|    August| 0.10192715231788088|
| MT|     January|0.010774703557312242|      May| 0.08077929906139672|
| NC|     January|0.042857036651550114|      July| 0.1053681637641307|
| VT|     January| 0.08598721023181453|      July| 0.24016064257028116|
| MD|    February|0.036847265221878214|      June| 0.10497577917923413|
| DE|    February| 0.04197385620915033|      June| 0.13578431372549024|
| MO|     January| 0.06203433476394853|      June| 0.20050105600466106|
| VI|     March| 0.04827044025157233|    October| 0.27268115942028986|
| IL|     January|0.030057494726815015|      May| 0.08716946191474498|
| ME|     March| 0.07291013950588417|      July| 0.16713952130764753|
| WA|       July|0.014176592672802735|    November| 0.18510913163515094|
| ND|    February|0.014421505376344075|      June| 0.10308791066095317|

```

MS	June	0.09098488888888889	December	0.17296009821976688
AL	June	0.09880618965407002	July	0.18179667881733497
IN	November	0.06111773801634633	June	0.12585707146426783
OH	January	0.07794820143884895	June	0.15337783646322375
TN	February	0.10388005301524189	December	0.18651829871414452
NM	November	0.006521143639787703	July	0.09166127292340888
IA	January	0.008268656716417905	June	0.0809310227658852
PA	February	0.07852021478205941	June	0.1609032953639901
SD	January	0.011879822305354207	June	0.11713599231138885
NY	February	0.07456230833565655	July	0.16723314606741563
TX	February	0.025009921011058464	July	0.09383419460271157
WV	February	0.062468800868149724	June	0.15884563532187407
GA	November	0.05424255670503665	August	0.11492106723155303
MA	March	0.10883711668657914	July	0.19126041666666663
KS	January	0.017501097935880543	June	0.19463163841807882
CO	February	0.0115337224383917	August	0.07849496184984579
FL	November	0.04829176678542824	July	0.26171101460718466
AK	March	0.028318786505874275	September	0.08177891420483052
AR	November	0.07980119423955044	April	0.1912900976290096
OK	January	0.013254607264269103	June	0.10921756437367833
PR	February	0.00864957264957265	May	0.2540740740740741
UT	July	0.02149922890504515	October	0.05086682697812897
HI	June	0.025426452410383178	December	0.16028352919231498

+-----+-----+-----+-----+-----+

Time to perform Task 3: 102.798s

Task 4

Finally, to perform Task 4, I needed the DataSet created in Task 3. I simply use `expr()` and took the difference between the highest precipitation column and the lowest precipitation column to create the Difference column. In addition, there is a built-in sort that can sort the rows in ascending order without any manipulation.

Please see below for the results for this task:

State	Lowest Month	Lowest Precipitation	Highest Month	Highest Precipitation	Difference
DC	January	0.0	January	0.0	0.0
NV	August	0.009235637779941571	July	0.027827986142491327	0.018592348362549756
UT	July	0.02149922890504515	October	0.05086682697812897	0.029367598073083822
RI	February	0.05318513451892387	April	0.08669625871908687	0.033511124200163
MI	February	0.03219802952210654	September	0.07008190283855044	0.0378838733164439
MN	January	0.003336326793773...	June	0.047394458223324575	0.04405813142955098
ID	August	0.01603681242685391	November	0.06329878524462283	0.04726197281776892
WY	February	0.011197420634920623	May	0.06364933741080532	0.0524519167758847
AK	March	0.028318786505874275	September	0.08177891420483052	0.05346012769895625
VA	February	0.024673968515683268	June	0.08022623892745996	0.0555522704117767
IL	January	0.030057494726815015	May	0.08716946191474498	0.05711196718792996
GA	November	0.05424255670503665	August	0.11492106723155303	0.06067851052651638
AZ	November	0.008935691105371218	July	0.07065813761059134	0.061722446505220116
NC	January	0.042857036651550114	July	0.1053681637641307	0.06251112711258058
WI	January	0.017938713700338508	June	0.08139457959880501	0.06345586589846651
IN	November	0.06111773801634633	June	0.12585707146426783	0.0647393334479215
CO	February	0.0115337224383917	August	0.07849496184984579	0.0669612394114541
MD	February	0.036847265221878214	June	0.10497577917923413	0.06812851395735592
TX	February	0.025009921011058464	July	0.09383419460271157	0.0688242735916531
MT	January	0.010774703557312242	May	0.08077929906139672	0.07000459550408447
SC	November	0.07339613086194678	August	0.14576830669591567	0.07237217583396889
IA	January	0.008268656716417905	June	0.0809310227658852	0.07266236604946728
OH	January	0.07794820143884895	June	0.15337783646322375	0.0754296350243748
CA	July	0.002754643909795044	February	0.08149602463781755	0.0787413807280225
KY	November	0.09035292213048407	July	0.16976356251813182	0.07941064038764775
CT	March	0.07905631114292533	June	0.1607084785133567	0.08165216737043138
MS	June	0.09098488888888889	December	0.17296009821976688	0.08197520933087799
PA	February	0.07852021478205941	June	0.1609032953639901	0.08238308058193068
MA	March	0.10883711668657914	July	0.19126041666666663	0.08242329998008749
TN	February	0.10388005301524189	December	0.18651829871414452	0.08263824569890263
LA	November	0.0869480834546337	July	0.1697039196631033	0.0827558362084696
AL	June	0.09880618965407002	July	0.18179667881733497	0.08299048916326494
NM	November	0.006521143639787703	July	0.09166127292340888	0.08514012928362118
ND	February	0.014421505376344075	June	0.10308791066095317	0.08866640528460909
NY	February	0.07456230833565655	July	0.16723314606741563	0.09267083773175908
DE	February	0.04197385620915033	June	0.13578431372549024	0.0938104575163399
ME	March	0.07291013950588417	July	0.16713952130764753	0.09422938180176336
NE	January	0.007690058479532156	August	0.10192715231788088	0.09423709383834873

OK	January	0.013254607264269103	June	0.10921756437367833	0.09596295710940922
WV	February	0.062468800868149724	June	0.15884563532187407	0.09637683445372434
NJ	February	0.07410393700787403	July	0.1785957240038873	0.10449178699601326
SD	January	0.011879822305354207	June	0.11713599231138885	0.10525617000603464
AR	November	0.07980119423955044	April	0.1912900976290096	0.11148890338945915
OR	July	0.00966687020829092	November	0.14075152653828096	0.13108465632999006
HI	June	0.025426452410383178	December	0.16028352919231498	0.1348570767819318
MO	January	0.06203433476394853	June	0.20050105600466106	0.13846672124071252
NH	March	0.08385728710871762	July	0.23177604593929466	0.14791875883057704
VT	January	0.08598721023181453	July	0.24016064257028116	0.15417343233846664
WA	July	0.014176592672802735	November	0.18510913163515094	0.1709325389623482
KS	January	0.017501097935880543	June	0.19463163841807882	0.17713054048219828
FL	November	0.04829176678542824	July	0.26171101460718466	0.21341924782175642
VI	March	0.04827044025157233	October	0.27268115942028986	0.22441071916871752
PR	February	0.00864957264957265	May	0.2540740740740741	0.24542450142450148

Time to perform Task 4: 104.624s